

# Hieu Minh “Jord” Nguyen

[jordnguyen43@gmail.com](mailto:jordnguyen43@gmail.com) | [Google Scholar](#) | [GitHub](#) | [LessWrong](#)

## Research Experience

---

### Center on Long-term Risk

Jul – Sep 2026

*Research Fellow*

- Studying LLM psychology and trait entanglement

### Cadenza Labs

Apr – Jun 2026

*Research Contractor*

- Building model organisms to benchmark lie detectors for LLMs
- Organised [Aletheia’s Quest](#) with [NDIF](#) and [Schmidt Sciences](#)

### Algoverse

Feb – Mar 2026

*Research Mentor*

- Mentored a research project on evaluation awareness.

### Constellation Institute

Jan – Feb 2026

*Visiting Researcher*

- Led research on investigating LLM introspective access and self-reports in collaboration with [Reciprocal Research](#)
- Continued research on activation monitoring and oversight

### Independent Contracting

Oct 2025 – Jan 2026

*Research Lead*

- Research on latent space monitors to improve AI oversight (e.g., probes for AI debate)
- Collaborating with researchers from PIBBSS, Stanford University, and University of Bristol
- Supported by a research grant from Open Philanthropy

### Trajectory Labs

Oct – Dec 2025

*Redteaming Contractor*

- Red-teamed AI agents for frontier AI companies
- Built RL environments for AI robustness training

### Pivotal Research

Feb – Apr 2025

*Technical Governance Research Fellow*

- Led research on [measuring and monitoring evaluation awareness](#)
- Used linear probes to detect when models recognise evaluation contexts
- Collaborated with researchers from Apollo Research, UCLA, and Waseda University

### Apart Research

Feb 2024 – Feb 2025

*Research Fellow*

- Led and published 3 research papers (ICLR, AAAI, TAIS)
- 1st place AI Security Evaluation Hackathon; 3rd place LLM Evaluations Hackathon

## Publications

---

|                                                                                                                                                                               |                                |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------|
| <b>Probing and Steering Evaluation Awareness of Language Models</b> <a href="#">[paper]</a><br>Technical AI Governance & Actionable Interpretability Workshops. First author. | <i>ICML 2025</i>               |
| <b>DarkBench: Benchmarking Dark Patterns in Large Language Models</b> <a href="#">[paper]</a><br>Oral Spotlight. Co-first author.                                             | <i>ICLR 2025</i>               |
| <b>Evaluating Precise Geolocation Inference Capabilities of VLMs</b> <a href="#">[paper]</a><br>Datasets and Evaluators of AI Safety Workshop. Co-first author.               | <i>AAAI 2025</i>               |
| <b>Identifying Cooperative Personalities through Personality Steering</b> <a href="#">[paper]</a><br>Co-author.                                                               | <i>TAIS 2025</i>               |
| <b>Evaluating Generalization Capabilities of LLM-Based Agents in Mixed-Motive Scenarios Using Concordia</b> <a href="#">[paper]</a><br>Co-author.                             | <i>NeurIPS 2025</i>            |
| <b>Governing Automated Strategic Intelligence</b> <a href="#">[paper]</a><br>Co-author.                                                                                       | <i>TAIS 2026</i>               |
| <b>Measuring the Assistant's Harmlessness Preferences on the User Turn</b> <a href="#">[paper]</a><br>First-author.                                                           | <i>Submitted to EMNLP 2026</i> |

## Technical Writing

---

Here's 18 Applications of Deception Probes (Aug 2025) [\[post\]](#)  
Can SAE steering reveal sandbagging? (Apr 2025) [\[post\]](#)  
A Survey of Theory of Mind in Large Language Models (Feb 2025, AAAI ToM4AI Workshop) [\[paper\]](#)  
Shallow Review of Technical AI Safety 2025 (Dec 2025) [\[post\]](#)  
Shallow Review of Technical AI Safety 2024 (Dec 2024) [\[post\]](#)

## Selected Talks

---

**NeurIPS 2024** – Concordia Workshop: Narrative steering for LLM cooperation [\[recording\]](#)  
**REAIM Summit 2024** – Invited by UN Office for Disarmament Affairs to present on catastrophic AI risks  
**Youth4Disarmament Forum 2025** – Invited to first-ever Y4D Forum at UN Headquarters to discuss risks from emerging technologies  
**ASEAN Regional Forum on Non-Proliferation and Disarmament 2025** – Nominated by UNODA, selected by co-chairs (Australia, NZ, Philippines) as a panelist  
**EA Summit Vietnam 2025** – Introduced AI safety research and catastrophic risks to the general public

## Selected Awards

---

1st prize, Sci-fAI Futures Youth Challenge (UNODA), 2024 – 8 winners from 230+ submissions  
1st place, AI Security Evaluation Hackathon (Apart Research), 2024  
3rd place, ML Model Attribution Challenge (DEFCON 30 AI Village), 2022 – presented at IEEE SaTML 2023

## Volunteering

---

**ICML Mechanistic Interpretability Workshop** – Paper reviewer, 2026  
**ICML Technical AI Governance Workshop** – Paper reviewer, 2026  
**Apart Research** – Judge, AI Control Hackathon and AI Manipulation Hackathon, 2026  
**Non-Trivial Fellowship** – Interviewer and facilitator across 2024 and 2025 cohorts.

**AI Safety Vietnam** – Founded and managed technical research at the [Hanoi AI Safety Network](#) and [AI Safety Vietnam](#)., 2023-2025

## Skills

---

**Interpretability:** TransformerLens, nnsight

**ML:** PyTorch, HuggingFace, LLM APIs (OpenAI, Anthropic, Google), LLM finetuning

**Infrastructure:** Git, Linux, remote compute (Runpod, vast.ai), running distributed experiment jobs

## Education

---

**University of Science and Technology of Hanoi** – BSc. Data Science (Class of 2026)

GPA: 3.7/4 | ACT: 36 | Academic Excellence Scholarship

**AI Security Bootcamp** (Aug 2025) – 4-week intensive on security fundamentals for AI systems.

Cryptography, network security, penetration testing, reverse engineering, adversarial ML, LLM security.

**Non-trivial Fellowship** (Sep – Dec 2023) – 48/4,800 applicants, 1% acceptance rate