

Hieu Minh “Jord” Nguyen

jordnguyen43@gmail.com | [Google Scholar](#) | [GitHub](#) | [LessWrong](#)

Research Experience

Center on Long-term Risk

Jul – Sep 2026

Research Fellow

- Studying LLM psychology and personas

Cadenza Labs

Apr – Jun 2026

Research Contractor

- Building model organisms to benchmark lie detectors for LLMs

Algoverse

Feb – Mar 2026

Research Mentor

- Mentored a research project on evaluation awareness.

Constellation Institute

Jan – Feb 2026

Visiting Researcher

- Led research on investigating LLM introspective access and self-reports in collaboration with [Reciprocal Research](#)
- Continued research on activation monitoring and oversight

Independent Contracting

Oct 2025 – Jan 2026

Research Lead

- Research on latent space monitors to improve AI oversight (e.g., probes for AI debate)
- Collaborating with researchers from PIBBSS, Stanford University, and University of Bristol
- Supported by a research grant from Open Philanthropy

Trajectory Labs

Oct – Dec 2025

Redteaming Contractor

- Red-teamed AI agents for frontier AI companies
- Built RL environments for AI robustness training

Pivotal Research

Feb – Apr 2025

Technical Governance Research Fellow

- Led research on [measuring and monitoring evaluation awareness](#)
- Used linear probes to detect when models recognise evaluation contexts
- Collaborated with researchers from Apollo Research, UCLA, and Waseda University

Apart Research

Feb 2024 – Feb 2025

Research Fellow

- Led and published 3 research papers (ICLR, AAI, TAIS)
- 1st place AI Security Evaluation Hackathon; 3rd place LLM Evaluations Hackathon

Publications

Probing and Steering Evaluation Awareness of Language Models [paper] Technical AI Governance & Actionable Interpretability Workshops. First author.	<i>ICML 2025</i>
DarkBench: Benchmarking Dark Patterns in Large Language Models [paper] Oral Spotlight. Co-first author.	<i>ICLR 2025</i>
Evaluating Precise Geolocation Inference Capabilities of VLMs [paper] Datasets and Evaluators of AI Safety Workshop. Co-first author.	<i>AAAI 2025</i>
Identifying Cooperative Personalities through Personality Steering [paper] Co-author.	<i>TAIS 2025</i>
Evaluating Generalization Capabilities of LLM-Based Agents in Mixed-Motive Scenarios Using Concordia [paper] Co-author.	<i>NeurIPS 2025</i>
Governing Automated Strategic Intelligence Co-author.	<i>TAIS 2026</i>

Technical Writing

Here's 18 Applications of Deception Probes (Aug 2025) [\[post\]](#)
Can SAE steering reveal sandbagging? (Apr 2025) [\[post\]](#)
A Survey of Theory of Mind in Large Language Models (Feb 2025, AAAI ToM4AI Workshop) [\[paper\]](#)
Shallow Review of Technical AI Safety 2025 (Dec 2025) [\[post\]](#)
Shallow Review of Technical AI Safety 2024 (Dec 2024) [\[post\]](#)

Selected Talks

NeurIPS 2024 – Concordia Workshop: Narrative steering for LLM cooperation [\[recording\]](#)
REAIM Summit 2024 – Invited by UN Office for Disarmament Affairs to present on catastrophic AI risks
Youth4Disarmament Forum 2025 – Invited to first-ever Y4D Forum at UN Headquarters, funded by UNODA
ASEAN Regional Forum on Non-Proliferation and Disarmament 2025 – Nominated by UNODA, selected by co-chairs (Australia, NZ, Philippines) as panelist
EA Summit Vietnam 2025 – Introduced AI safety research and catastrophic risks to the public

Selected Awards

1st prize, Sci-fAI Futures Youth Challenge (UNODA), 2024 – 8 winners from 230+ submissions
1st place, AI Security Evaluation Hackathon (Apart Research), 2024
3rd place, ML Model Attribution Challenge (DEFCON 30 AI Village), 2022 – presented at IEEE SaTML 2023

Volunteering

ICML Technical AI Governance Workshop – Paper reviewer, 2026
Apart Research – Judge, AI Control Hackathon and AI Manipulation Hackathon, 2026
Non-Trivial Fellowship – Interviewer and facilitator across 2024 and 2025 cohorts.

Skills

Interpretability: TransformerLens, nnsight
ML: PyTorch, HuggingFace, LLM APIs (OpenAI, Anthropic, Google), LLM finetuning
Infrastructure: Git, Linux, remote compute (Runpod, vast.ai), running distributed experiment jobs

Education

University of Science and Technology of Hanoi – BSc. Data Science (Class of 2026)

GPA: 3.7/4 | ACT: 36 | Academic Excellence Scholarship

AI Security Bootcamp (Aug 2025) – 4-week intensive on security fundamentals for AI systems.

Cryptography, network security, penetration testing, reverse engineering, adversarial ML, LLM security.

Non-trivial Fellowship (Sep – Dec 2023) – 48/4,800 applicants, 1% acceptance rate